

Response to Review of paper: Harnessing Frequency-Aware Network for Crack Detection in Road-Perceptive Autonomous Vehicles Scene

Senyun Kuang, Yintao Wei,
Yang Liu, and Xiaobo Qu, *Senior Member, IEEE*

Dear Associate Editor and Reviewers:

We sincerely thank you for the constructive feedback on our revised manuscript “Harnessing Frequency-Aware Network for Crack Detection in Road-Perceptive Autonomous Vehicles Scene” (CYB-E-2025-10-3956). We are encouraged that Reviewer 2 is fully satisfied with the previous revision and recommends acceptance, and we are grateful to the Associate Editor and Reviewer 1 for the remaining suggestions, all of which concern the clarity and completeness of the presentation.

In this second-round revision, we have carefully addressed every remaining comment. Specifically, we have (i) clarified the representativeness and selection criteria of all qualitative visualizations and added per-sample quantitative metrics to the corresponding captions; (ii) restructured the descriptions of the proposed modules (FAM, T-M, and HFM) in Sections III-B and III-C into an “intuition-first” style, providing a concise high-level explanation of each module before the mathematical formulation and consolidating the previously scattered equation-by-equation commentary; (iii) added a new subsection “Failure Cases and Limitations” (Section IV-H) that presents representative failure examples and analyzes the boundary conditions of the proposed method; and (iv) updated the Related Work section with recent (2024–2026) publications from journals and conferences, including papers from IEEE Transactions on Cybernetics.

For clarity of marking: all text that is newly added or modified *in this round* is highlighted in red in the revised manuscript, while the red markings of the previous round have been restored to black.

Below, we address each comment from the Associate Editor and the reviewers in turn. The comments are labelled in bold and blue, followed by our responses.

CONTENTS

I	Response to Associate Editor	3
II	Response to Reviewer 1	5
III	Response to Reviewer 2	10

I. RESPONSE TO ASSOCIATE EDITOR

Comments to the Author:

The revised manuscript has been substantially improved; however, further clarification is still needed regarding the representativeness of the visualizations, presentation of the proposed modules, and discussion of failure cases and limitations. I therefore recommend revision before acceptance. To put your contributions into the context of the most recent research, please update your Related Work section to include relevant references that have appeared in the past two years in conferences and journals, including from this Transaction.

Response: We sincerely appreciate your positive assessment of the previous revision and your time in handling our manuscript. All the remaining concerns have been fully addressed in this round:

- **Representativeness of the visualizations.** We now explicitly state the selection criteria of every qualitative figure, label the exact configuration used in each frequency-based visualization, and report the per-sample quantitative metric (IoU) of each visualized example directly in the caption, so that the qualitative and quantitative evidence are connected. Please see our detailed response to **Reviewer 1, Comment 1**.
- **Presentation of the proposed modules.** Sections III-B and III-C have been restructured in an “intuition-first” style: each module (FAM, T-M, and HFM) now begins with a concise, high-level explanation of its purpose and design rationale before any mathematical expression, and the previously scattered equation-by-equation commentary has been consolidated and condensed. Please see our detailed response to **Reviewer 1, Comment 2**.
- **Failure cases and limitations.** A new subsection, “Failure Cases and Limitations” (Section IV-H), has been added. It presents representative failure examples grouped into three dominant failure modes, analyzes why FANet struggles in these conditions, and discusses the practical boundary conditions and future directions. Please see our detailed response to **Reviewer 1, Comment 3**.
- **Related Work update.** We have updated the Related Work section with recent publications that appeared in the past two years in journals and conferences, including works from IEEE Transactions on Cybernetics. The newly added text and references are listed below.

In the revised manuscript, the following paragraph has been added to Section II (Related Work):

More recent studies from the past two years further advance both crack detection and frequency-aware representation learning. On the crack-detection side, UP-CrackNet [1] formulates road crack detection as unsupervised adversarial image restoration to remove the dependence on pixel-wise annotations; SCSEgamba [2] designs a lightweight structure-aware vision Mamba that captures crack morphology and texture with minimal computational cost; and an ultra-lightweight Mamba network with frequency-domain feature extraction has been developed for pavement tiny-crack segmentation [3]. On the frequency-modeling side, FreqFusion [4] rectifies feature fusion for dense prediction with adaptive low-/high-pass filters, and FSEL [5] entangles frequency and spatial representations for camouflaged object detection. Closely related inspection and segmentation problems have also been actively studied in IEEE Transactions on Cybernetics, including zero-shot industrial anomaly detection with hybrid vision-language prompts [6] and contrastive transformer-based segmentation of subtle structures [7].

These advances confirm the growing interest in spectral cues and efficient architectures; however, they either address generic dense prediction or treat frequency information as an auxiliary enhancement, leaving a deeply coupled frequency-spatial representation dedicated to road crack detection, the very focus of FANet, largely unexplored.

The above paragraph has been added to the end of **Section II (Related Work)**, and the seven newly cited works are: two from IEEE Transactions on Cybernetics (AnomalyVLM, TCYB 2025 [6]; CTNet, TCYB 2024 [7]), three recent crack-detection studies (UP-CrackNet, T-ITS 2024 [1]; SC-Segamba, CVPR 2025 [2]; ultra-lightweight Mamba–frequency network, ESWA 2025 [3]), and two recent frequency-aware representation studies (FreqFusion, TPAMI 2024 [4]; FSEL, ECCV 2024 [5]).

II. RESPONSE TO REVIEWER 1

General Comments:

In this paper, the authors propose FANet, a frequency–spatial synergistic framework for crack detection that integrates a Frequency-Aware Module (FAM), a Tri-Merge decoder (T-M), and a Hybrid Fusion Module (HFM). The revised manuscript provides expanded technical explanations, ablation studies involving low-/high-/full-frequency variants of FAM, and frequency-based visual comparisons. The work addresses a meaningful problem in pavement crack analysis and shows clear improvements over baseline models. However, several issues still require clarification or refinement before the paper can be accepted.

Response: We sincerely thank the reviewer for the careful reading of our revised manuscript, the accurate summary of our contributions, and the constructive suggestions. All three remaining comments have been fully addressed in this round, as detailed below.

Q1 Comments to the Author:

The manuscript includes frequency-based visualizations (high/low/fusion), but the representativeness and selection criteria of these examples should be clarified. The authors should explicitly state whether the provided samples are typical, randomly selected, or chosen as illustrative cases, and each visualization should clearly label the specific configuration used. Adding the corresponding quantitative metric of each visualized sample to the caption would help readers connect qualitative and quantitative evidence more directly.

Response: We thank the reviewer for this valuable suggestion. In the revised manuscript, we have made three changes to clarify the representativeness of the frequency-based visualizations:

- 1) **Selection criteria stated explicitly.** We now explicitly state in Section IV-G that the sample shown in the frequency-based visualization (Fig. 6) is a *typical case chosen as an illustrative example* (not randomly sampled, and not a best-performing case): it simultaneously contains a wide trunk crack and hairline branches on coarsely textured pavement, which is the condition under which the differences among the three FAM configurations are most visible. A corresponding statement of the selection criteria has also been added for the qualitative comparison figure (Fig. 5) in Section IV-F.
- 2) **Configuration of each panel labelled.** The caption of Fig. 6 now explicitly maps each panel to the exact configuration used in the ablation study (Table II): panel (c) is produced by *FAM (Full)*, panel (d) by *FAM (High-Freq only)*, and panel (e) by *FAM (Low-Freq only)*.
- 3) **Per-sample metrics added to the caption.** The per-sample IoU of each visualized prediction is now reported in the caption of Fig. 6, directly connecting the qualitative visualization with the quantitative evidence in Table II.

In the revised manuscript, the following text has been added to Section IV-G (Ablation Studies):

To clarify the representativeness of the frequency-based visualization, the sample in Fig. 6 is a typical case selected from the test set as an illustrative example, rather than a randomly drawn or best-

performing image: it simultaneously contains a wide trunk crack and hairline branches on coarsely textured pavement, the condition under which the differences among the three FAM configurations in Table II are most visible. Panels (c)–(e) are produced by the *FAM (Full)*, *FAM (High-Freq)*, and *FAM (Low-Freq)* configurations of Table II, respectively, and the per-sample IoU of each panel is reported in the caption to connect the qualitative and quantitative evidence.

The caption of Fig. 6 has been rewritten as follows (the figure is reproduced below; figure/table numbers in the quoted text refer to the revised manuscript):




Fig. 1. Qualitative comparison of the three FAM configurations in Table II on a typical sample selected from the test set as an illustrative case. (a) Input. (b) Ground truth. (c) Prediction of *FAM (Full)*, which fuses the high- and low-frequency branches, per-sample IoU = **0.79**. (d) Prediction of *FAM (High-Freq only)*, per-sample IoU = **0.61**. (e) Prediction of *FAM (Low-Freq only)*, per-sample IoU = **0.57**. The full-fusion configuration preserves both the main crack contour and the hairline branches, which is consistent with the quantitative gains reported in Table II.

In addition, the following sentence has been added to Section IV-F (Qualitative Results) to state the selection criteria of the method-comparison figure (Fig. 5):

The examples in Fig. 5 are typical cases selected from the test sets of the five datasets to cover the dominant challenge conditions (thin hairline cracks, coarse pavement texture, shadows and uneven illumination, and crack-like distractors) rather than randomly sampled or best-performing images; all methods are compared on identical inputs.

Q2 Comments to the Author:

The overall descriptions of FAM, T-M, and HFM—particularly in Sections III-B and III-C—remain verbose and would benefit from a clearer separation between intuition and formalism. Providing a concise, high-level explanation of each module’s purpose before introducing the mathematical expressions would significantly improve readability and help interdisciplinary readers better understand the design motivations.

Response: We thank the reviewer for this insightful suggestion, which has significantly improved the readability of the method section. In the revised manuscript, Sections III-B and III-C have been restructured following exactly the principle suggested by the reviewer, namely a clear separation between intuition and formalism:

- 1) **Intuition first.** Each of the three modules (FAM, T-M, and HFM) now begins with a concise “*Design intuition*” paragraph that explains, in plain language and without any equation, what problem the module solves and why it is designed this way.
- 2) **Formalism afterwards, condensed.** The mathematical formulation follows the intuition paragraph. The previously scattered equation-by-equation commentary (several separate explanatory

passages inserted after Eqs. (1)–(4) in the last revision) has been consolidated into brief remarks, and overlapping explanations of the FAM mechanism, which appeared three times in slightly different wording, have been merged into a single passage. As a result, Sections III-B and III-C are now *shorter* than in the previous version while conveying the same technical content.

In the revised manuscript, the newly added “Design intuition” paragraphs are shown below:

(Section III-B, before the formulation of FAM) Design intuition of FAM. Road cracks appear as sparse and sharp intensity discontinuities embedded in slowly varying pavement context. In the feature space, the former concentrates in the high-frequency band, whereas the latter dominates the low-frequency band. The FAM is therefore designed to (i) decompose each backbone feature into a high-frequency and a low-frequency component, (ii) let the two components exchange complementary information, so that edge responses acquire contextual support and contour responses are cleansed of texture noise, and (iii) adaptively re-weight and merge the two bands with learnable gates. The output is a frequency-balanced feature in which crack evidence is amplified relative to the background texture. The remainder of this subsection formalizes these three steps.

(Section III-B, before the formulation of the T-M decoder) Design intuition of T-M. The T-M decoder aggregates the frequency-aware features of the three deepest levels: deeper features indicate *where* a crack is likely to exist, while shallower ones describe its shape more precisely. Multiplying adjacent levels retains only the responses that are confirmed by both, yielding a clean coarse localization map S_1 .

(Section III-C, before the formulation of HFM) Design intuition of HFM. The HFM addresses two questions during refinement. The first is *where to refine*: the coarse map S_g acts as a spatial prior that amplifies crack-relevant regions of the deeper feature. The second is *what to keep*: a channel-level association between adjacent layers determines which high-resolution details are consistent with this prior and should be preserved. Both operations are lightweight, so the refinement improves accuracy at a marginal computational cost. The following formulation details the two steps.

In addition, the redundant explanatory passages of the previous round have been merged and condensed (the removed redundancies are not marked in red, to keep the marking readable). We believe the method section is now significantly easier to follow, especially for interdisciplinary readers.

Q3 Comments to the Author:

The paper would benefit from a short subsection discussing typical failure cases or limitations. Presenting a small number of representative failure examples and briefly analyzing why FANet struggles in those conditions would provide a more balanced understanding of the model’s behavior and practical boundary conditions.

Response: We thank the reviewer for this constructive suggestion, which makes the experimental analysis more balanced. In the revised manuscript, we have added a new subsection, “**Failure Cases and Limitations**” (Section IV-H), placed after the ablation studies. It presents representative failure examples, taken directly from FANet’s own predictions on the hardest crack categories, grouped into

two dominant failure modes, analyzes the underlying reasons within the frequency-aware design, and summarizes the practical boundary conditions and future directions of FANet.

In the revised manuscript, the newly added subsection is shown below (figure numbers in the quoted text refer to the revised manuscript, where the failure-case figure is Fig. 7; it is reproduced below):

IV-H. Failure Cases and Limitations. Although FANet achieves state-of-the-art performance on all five datasets, it still struggles under two identifiable conditions. Fig. 7 shows representative failure cases taken from FANet’s own predictions on the most challenging crack categories. In each column, the top row is the input image, the middle row is the ground truth, and the bottom row is our prediction, so that the gap between the ground truth and our result can be read directly.

(1) Fragmentation of low-contrast and camouflaged cracks. When a crack blends into the pavement and its intensity contrast is low, the high-frequency response of FAM becomes weak and intermittent, while the low-frequency contour prior is too smooth to bridge the gaps. As shown in Fig. 7(a), the crack that forms a connected structure in the ground truth is broken into disconnected segments in our prediction.

(2) Under-segmentation of fine hairline cracks. When cracks are only a few pixels wide, the frequency evidence they produce is faint relative to the surrounding pavement texture. As shown in Fig. 7(b) and (c), the thin cracks that are complete in the ground truth are only weakly activated in our prediction, and some fine branches are missed.

These two modes delineate the practical boundary condition of the proposed design. FANet implicitly assumes that cracks constitute the dominant sparse high-frequency anomaly of the scene, an assumption that is violated when the contrast is very low or the crack is extremely thin. In addition, two limitations are worth noting. First, the two-step coarse-to-fine pipeline introduces extra inference cost compared with single-pass lightweight detectors, as reflected by the FPS comparison in Table I. Second, FANet is trained with the binary cross-entropy loss only and imposes no explicit topological continuity constraint, which partially explains the fragmented predictions above. We plan to explore topology-aware objectives, contrast-adaptive frequency gating, and complementary sensing modalities such as depth or infrared imaging to push these boundaries in future work.




Fig. 2. Representative failure cases of FANet on the most challenging crack categories. In each column, the top row is the input image, the middle row is the ground truth, and the bottom row is our prediction. (a) A camouflaged, low-contrast crack whose connected ground-truth structure is fragmented into disconnected segments by our prediction. (b), (c) Fine hairline cracks whose complete ground truth is only weakly and partially recovered by our prediction.

III. RESPONSE TO REVIEWER 2

Comments to the Author:

The authors have addressed all my concerns. I recommend acceptance. Thanks.

Response: We sincerely thank the reviewer for the positive evaluation of our revised manuscript and for recommending acceptance. The reviewer's constructive comments in the previous round have contributed greatly to improving the quality of this work.

REFERENCES

- [1] N. Ma, R. Fan, and L. Xie, “UP-CrackNet: unsupervised pixel-wise road crack detection via adversarial image restoration,” *IEEE Trans. Intell. Transp. Syst.*, vol. 25, pp. 13 926–13 936, 2024.
- [2] H. Liu, C. Jia, F. Shi, X. Cheng, and S. Chen, “SCSegamba: lightweight structure-aware vision mamba for crack segmentation in structures,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2025*, 2025, pp. 29 406–29 416.
- [3] Z. Zhang, B. Peng, and T. Zhao, “An ultra-lightweight network combining Mamba and frequency-domain feature extraction for pavement tiny-crack segmentation,” *Expert Syst. Appl.*, vol. 264, p. 125941, 2025.
- [4] L. Chen, Y. Fu, L. Gu, C. Yan, T. Harada, and G. Huang, “Frequency-aware feature fusion for dense image prediction,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 46, no. 12, pp. 10 763–10 780, 2024.
- [5] Y. Sun, C. Xu, J. Yang, H. Xuan, and L. Luo, “Frequency-spatial entanglement learning for camouflaged object detection,” in *European Conference on Computer Vision, ECCV 2024*, 2024.
- [6] Y. Cao, X. Xu, C. Sun, Y. Cheng, Z. Du, L. Gao, and W. Shen, “Personalizing vision-language models with hybrid prompts for zero-shot anomaly detection,” *IEEE Trans. Cybern.*, vol. 55, no. 4, pp. 1917–1929, 2025.
- [7] B. Xiao, J. Hu, W. Li, C.-M. Pun, and X. Bi, “CTNet: contrastive transformer network for polyp segmentation,” *IEEE Trans. Cybern.*, vol. 54, no. 9, pp. 5040–5053, 2024.